

Exploring IBM watsonx Prompt Lab

Estimated time: 40 minutes

When it comes to prompting foundation models, there isn't just one right answer. There are usually multiple ways to prompt a foundation model for a successful result.

The **IBM watsonx Prompt Lab** is a prompting tool designed to help you craft prompts that achieve the desired or expected result from a given model and set of parameters. In this activity, you will explore the key features of the watsonx Prompt Lab, and also learn how to prompt a foundation model.

Accessing IBM watsonx

To access the IBM watsonx platform, you need to have an IBM cloud account. You can use the credentials for the IBM cloud account to access the watsonx platform.

If you do not have an IBM cloud, you will be asked to register and create your account when you try to access the IBM watsonx platform.

Here is the link to the platform: [IBM watsonx](#).

2. Click the user icon on the upper right corner of the IBM watsonx.ai home screen to log in.

Meet watsonx

The AI and data platform that's built for business

Start your free trial →

watsonx.ai

Train, validate, tune and deploy models for generative AI

Start building AI applications →

watsonx.data

Bring all your business data together to scale analytics and AI in your applications

Manage your data for AI →

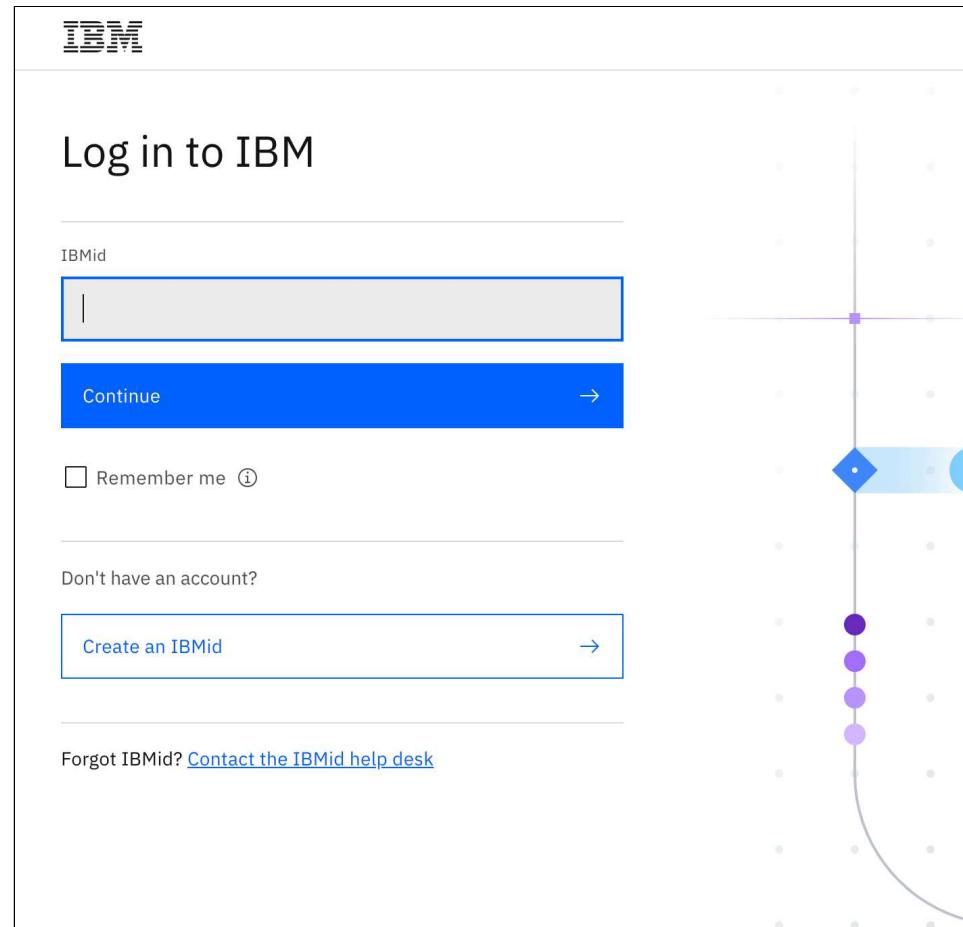
watsonx.governance

Manage, monitor and govern all your AI

watsonx Assistant

Get started with a pre-built application for

3. Log in using the credentials you used while signing up with the feature code.



The image shows the IBM login page. At the top left is the IBM logo. Below it is the heading "Log in to IBM". A text input field labeled "IBMid" is followed by a blue "Continue" button with a white arrow. To the left of the "Remember me" checkbox is a small informational icon. Below the login area is a link "Create an IBMid" with a blue arrow. At the bottom is a link "Forgot IBMid? [Contact the IBMid help desk](#)". On the right side of the page, there is a decorative graphic consisting of a vertical grey line with several small grey dots. Along this line are several colored shapes: a purple square at the top, a blue diamond below it, a light blue rectangle, and a blue circle further down. To the right of these shapes is a curved grey line.

IBMid

Continue →

Remember me ⓘ

Don't have an account?

Create an IBMid →

Forgot IBMid? [Contact the IBMid help desk](#)

Note: You might receive a verification code via email to verify and login to your IBM watsonx account.

4. After logging in, you will land on the home page, where you need to click 'Start your free trial' and log in again.

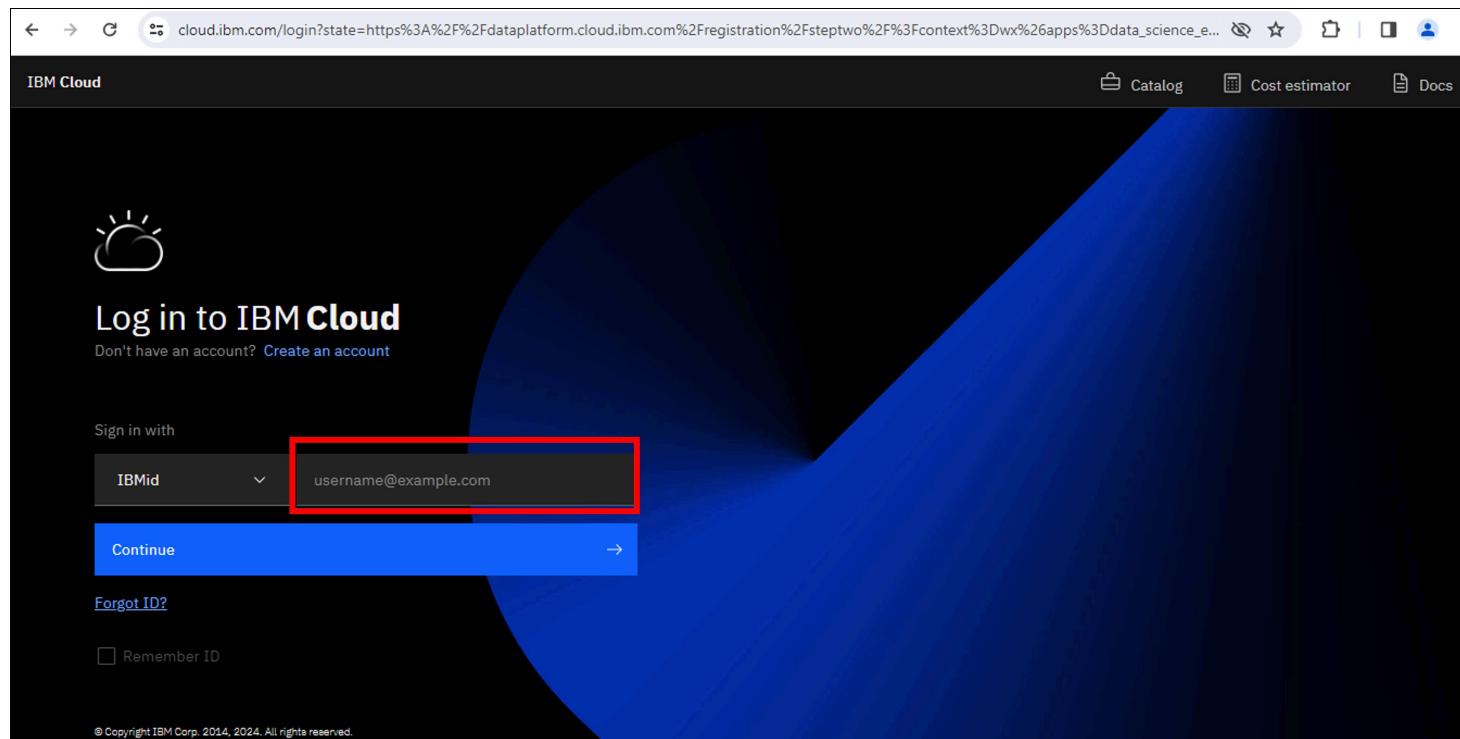
The screenshot shows the IBM WatsonX homepage. At the top, there's a navigation bar with links for Products, Solutions, Consulting, Support, Think, WatsonX (which is underlined), AI platform, Models, AI assistants, Technology partners, Resources, Pricing, Book a live demo, and Try it for free. Below the navigation, a large banner features the text "Meet watsonx" and "The AI and data platform that's built for business". A prominent red button labeled "Start your free trial" with an arrow points to the right. To the right of the banner, there are four main sections: "watsonx.ai" (with a grid icon) which says "Train, validate, tune and deploy models for generative AI"; "watsonx.data" (with a hexagon icon) which says "Bring all your business data together to scale analytics and AI in your applications"; "watsonx.governance" (with a scales icon) which says "Manage, monitor and govern all your AI"; and "watsonx Assistant" (with a person icon) which says "Get started with a pre-built application for".

5. Click on **Create account** or **Log in**. In the next step, click on **Log in**, along the text Already have an IBM Cloud account. Choose the same region you had used to create your account.

The screenshot shows the IBM WatsonX sign-up page. The URL in the address bar is dataplatform.cloud.ibm.com/registration/stepone?context=wx&preselect_region=true. The page has a dark header with "IBM WatsonX" and a sub-header "Try IBM watsonx.ai for free". It says "Build, train, validate, tune and deploy AI models." Below this, there's a "Select a region" section with a dropdown set to "Dallas (us-south)". A note states: "Note: tools for foundation models are available only in the Dallas and Frankfurt regions." At the bottom, there's a "Create an IBM Cloud account" section with a checkbox for agreeing to terms and Data Use Policy, and a "Create account or log in" button. To the right, a mobile device screen displays the WatsonX interface with a "Welcome, Kate" message, workspace options like "Experiment with foundation models and build prompts", and project links for "flights-data_shaped.csv" and "Library 1 / Sample data for bank marketing".

The screenshot shows a web browser window with the URL cloud.ibm.com/registration?target=https%3A%2F%2Fdataplatform.cloud.ibm.com%2Fregistration%2Fsteptwo%2Fcontext%3Dwx%26apps%3Ddata_sci.... The page is titled "IBM Cloud" and features a "Create an IBM Cloud account" section. It includes a "Log in" button, which is highlighted with a red box. To the right, there is a large promotional banner for "watsonx.ai" with the tagline "Get started for free with watsonx.ai™". The banner describes "Build and deploy generative AI and machine learning solutions." and lists benefits such as working with open source or proprietary foundation models, using tools like Rstudio and Jupyter Notebooks, and managing the entire lifecycle from model development to management. A "Cookie Preferences" link is visible on the right side of the banner.

6. Log in again using the credentials you used while signing up with the feature code.



7. Provide the required information and then click on **Continue**.

IBM watsonx

Provide your information to continue

Company name

Phone number

+91 ▾

Continue

IBM may use my contact data to keep me informed of products, services, and offerings:

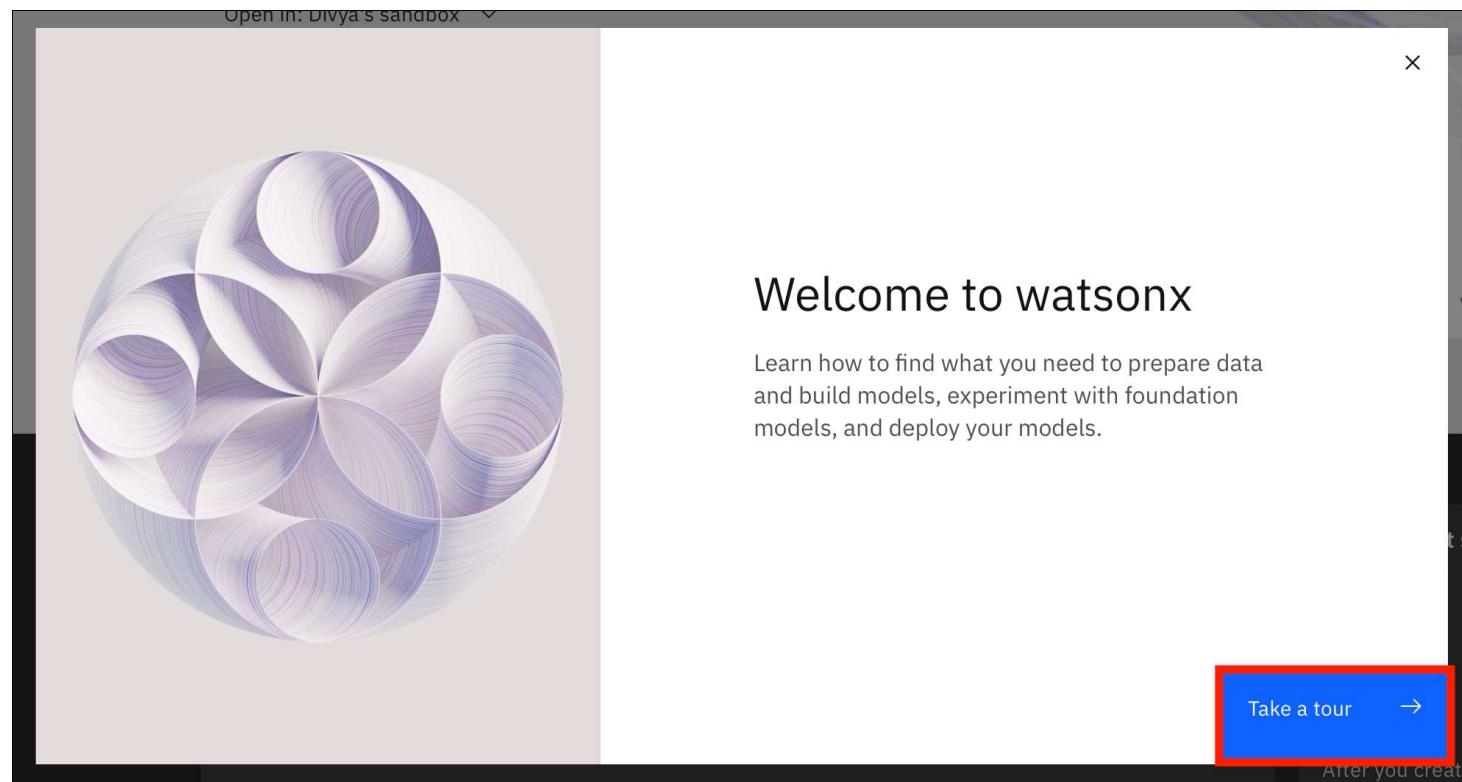
by email

by phone

You can withdraw your marketing consent at any time by submitting an [opt-out request](#). Also, you may unsubscribe from receiving marketing emails by clicking the unsubscribe link in each email. More information on our processing can be found in the [IBM Privacy Statement](#).

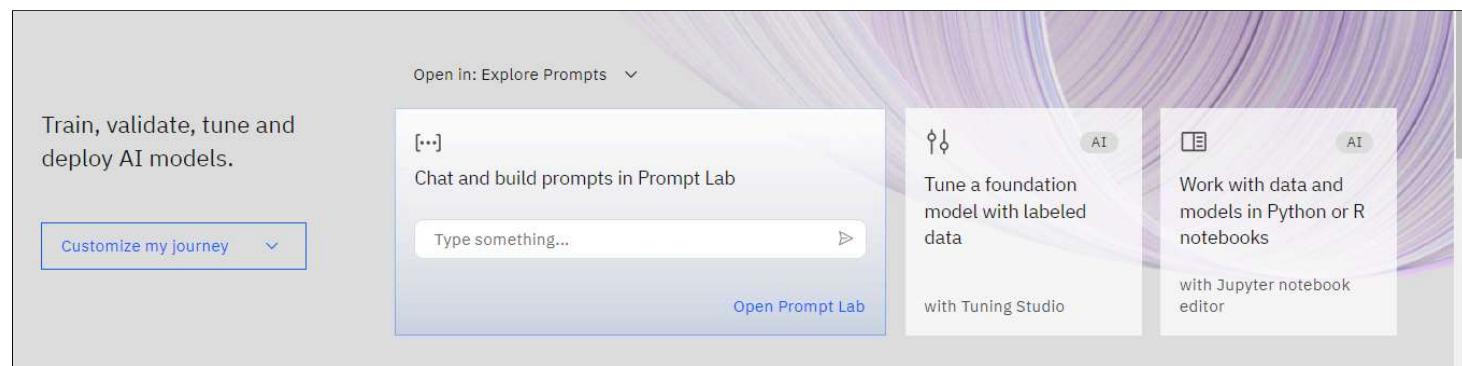
By submitting this form, I acknowledge that I have read and understand the IBM Privacy Statement and I accept the product [Terms and Conditions](#) of this registration form.

8. After the successful login, you will be prompted to **Take a tour** of the watsonx platform (on your first login), as shown below.



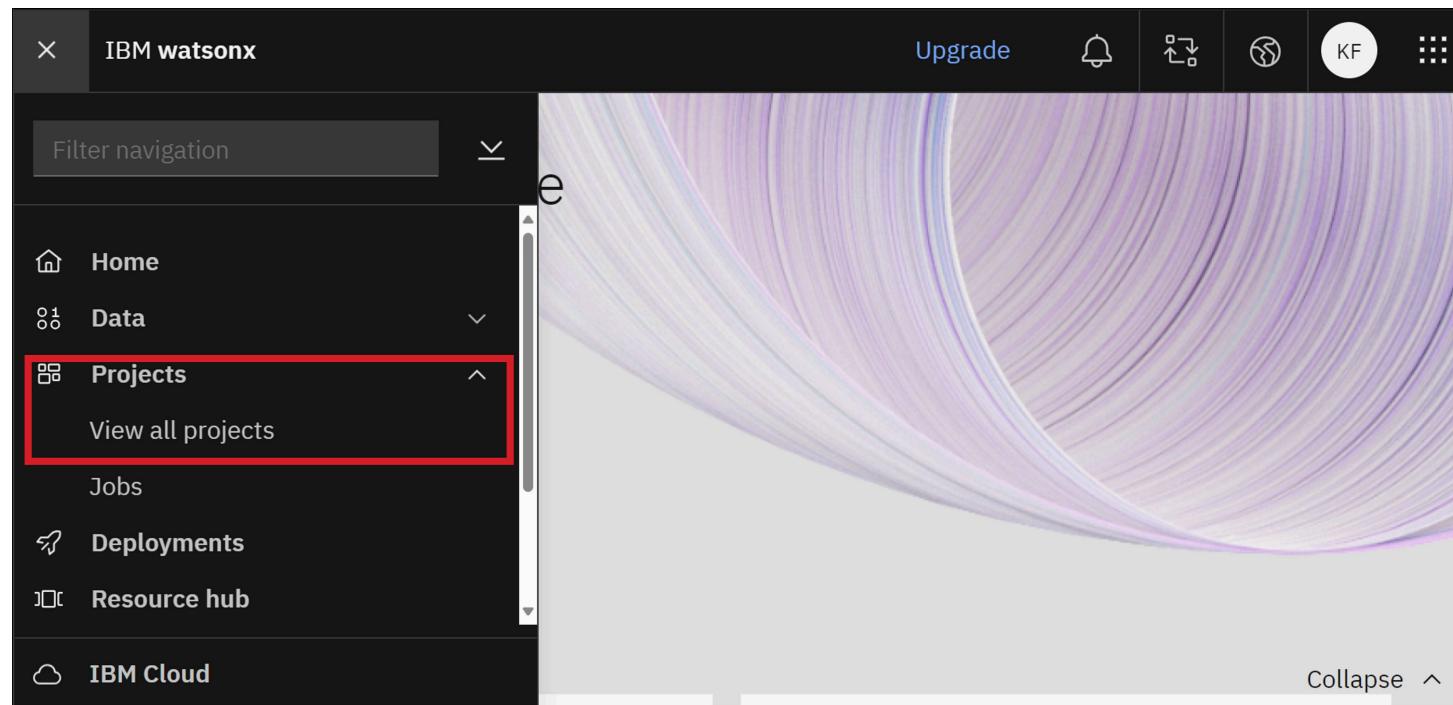
Note: You may opt for the watsonx walkthrough by clicking **Take a tour** or skip or exit it by clicking **X** at the upper right corner of the screen

- On completing or exiting the tour, you will reach your IBM watsonx.ai dashboard page, as shown below. Here, you can perform experiments like train, validate, tune, and deploy AI models.



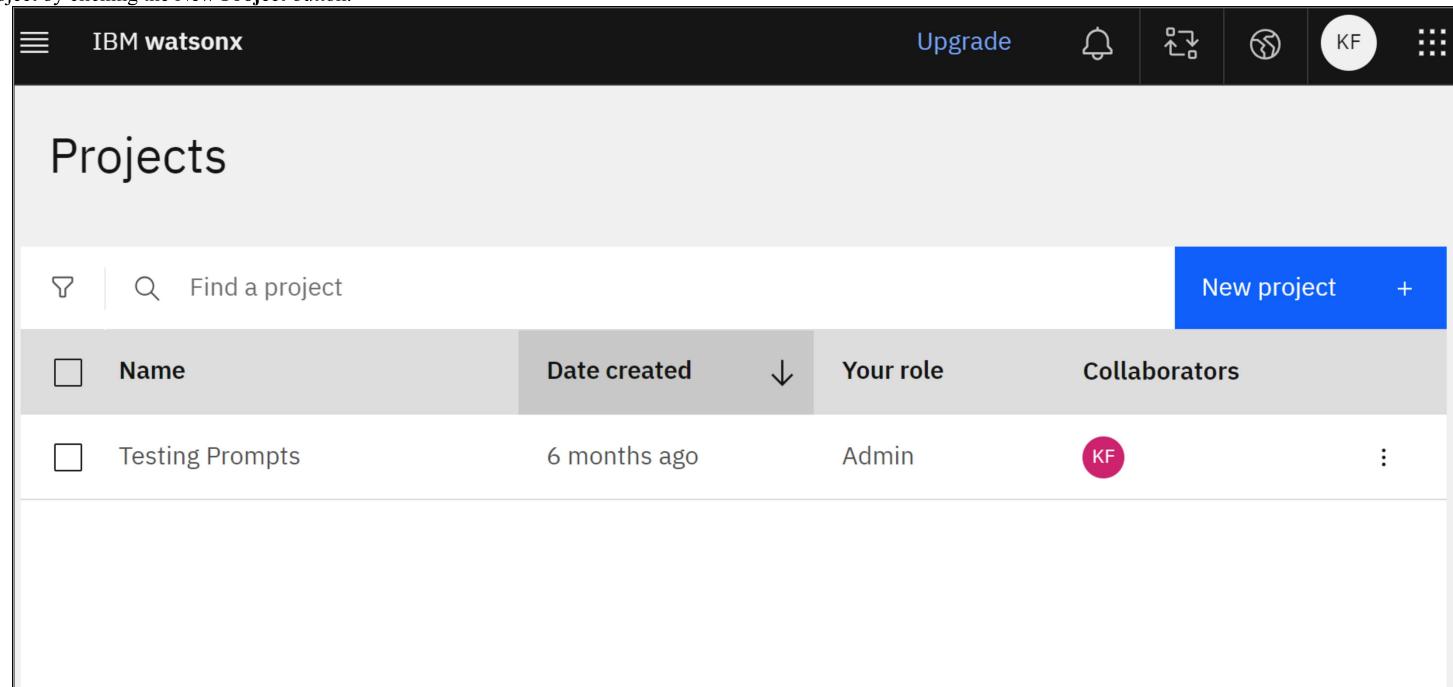
Note: If this is NOT your first login session in the watsonx.ai, you may find the **Prompt Lab** deactivated. In this case, you need to create a project and then select the **Open Prompt Lab** option. You can create a new project by following the given steps.

- You can view the list of projects for your account using the **Projects** drop-down option in the left panel.



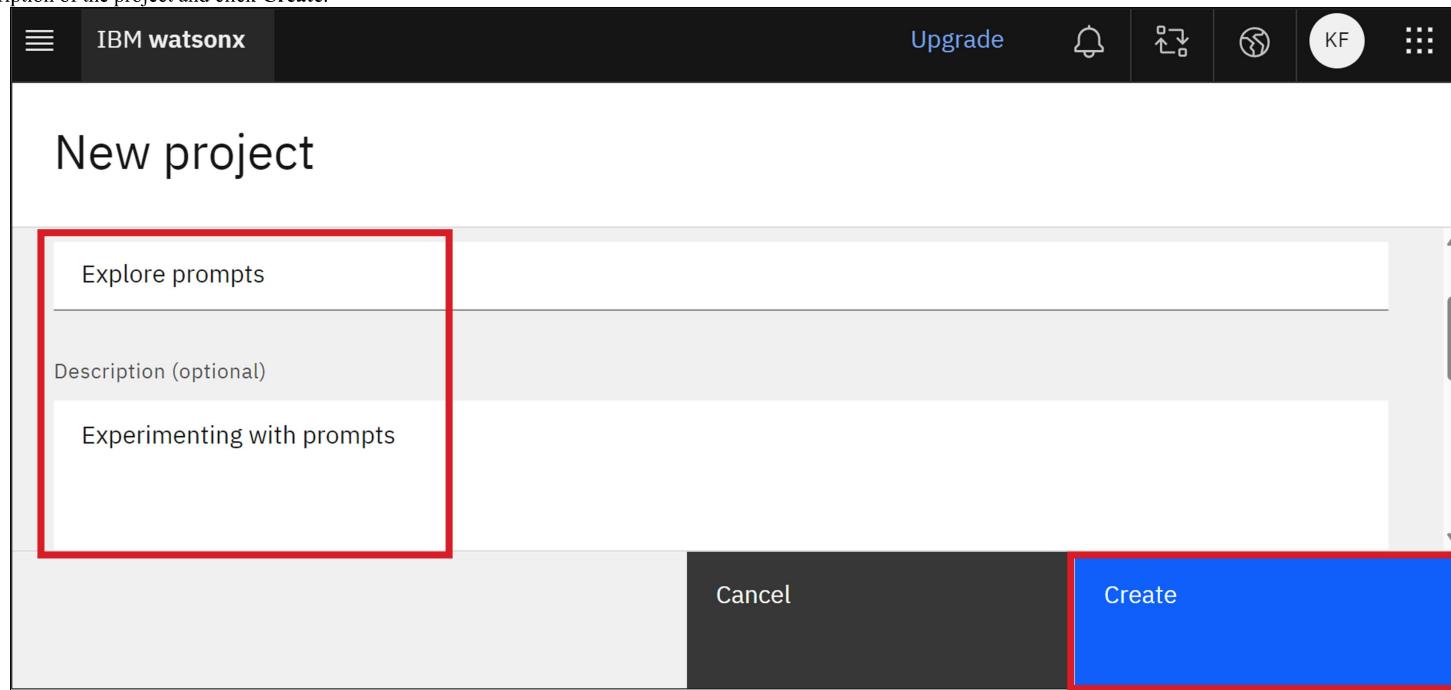
The screenshot shows the IBM WatsonX interface. On the left is a dark sidebar with white text and icons. At the top of the sidebar is a 'Filter navigation' dropdown. Below it are several menu items: 'Home', 'Data' (with a dropdown arrow), 'Projects' (which is highlighted with a red box), 'View all projects', 'Jobs', 'Deployments', 'Resource hub', and 'IBM Cloud'. To the right of the sidebar is a large, abstract purple and white curved graphic. At the top of the main area are several icons: a bell, a gear, a person icon, and a KF button. The URL 'about:blank' is visible at the very top center.

b. You can create a new project by clicking the **New Project** button.



The screenshot shows the 'Projects' page within the IBM WatsonX interface. The title 'Projects' is at the top. Below it is a search bar with a magnifying glass icon and the placeholder 'Find a project'. To the right of the search bar is a blue button labeled 'New project' with a '+' sign. The main area displays a table with one row of data. The columns are: a checkbox, 'Name' (containing 'Testing Prompts'), 'Date created' (containing '6 months ago'), 'Your role' (containing 'Admin'), 'Collaborators' (containing a circular icon with 'KF'), and a three-dot menu icon. The URL 'about:blank' is visible at the very top center.

c. Enter the name and description of the project and click **Create**.



d. Go to **Home** using drop-down option in the left panel.

e. You will reach your IBM watsonx.ai dashboard page. Click the first card with Prompt Lab, as shown in the screenshot below.



f. Click on **Associate Service**

No Watson Machine Learning service detected

To access the Prompt Lab, you must associate a Watson Machine Learning service instance to this project.

Associate service

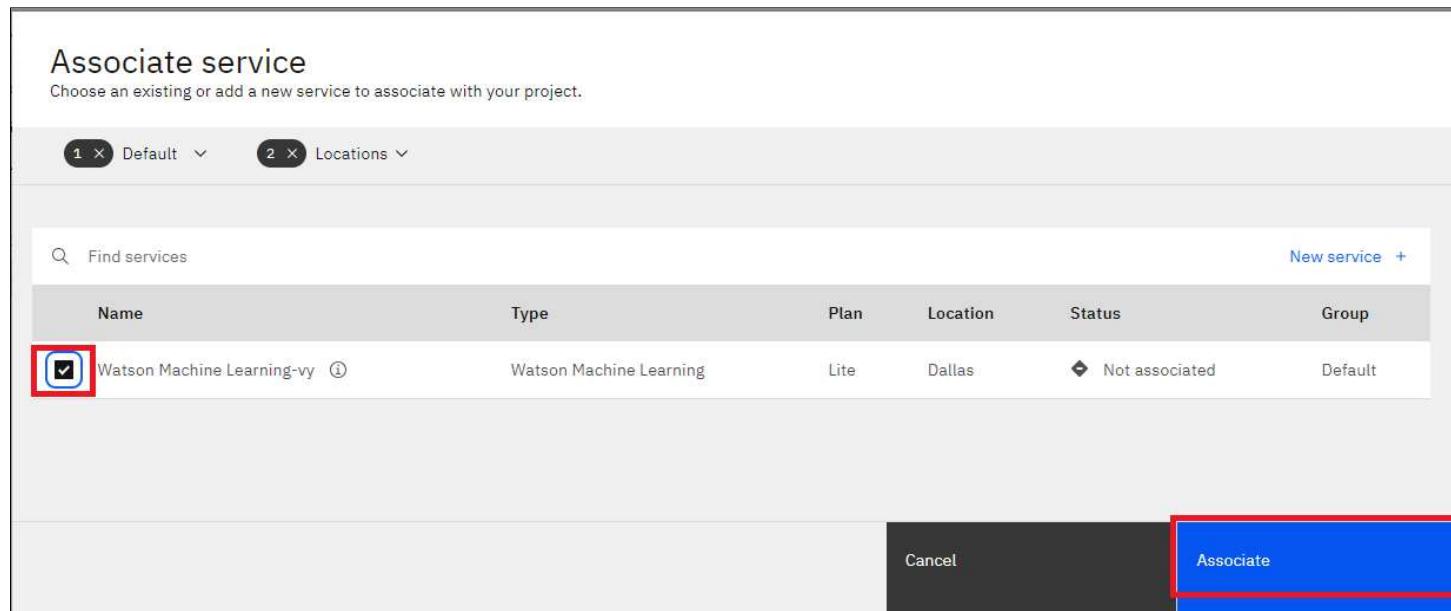
g. Select **Watson Machine Learning-vy** and click **Associate**. If Watson Machine Learning is not seen please select **New Service** and add **Watson Machine Learning**

Associate service
Choose an existing or add a new service to associate with your project.

1 Default 2 Locations

Name	Type	Plan	Location	Status	Group
<input checked="" type="checkbox"/> Watson Machine Learning-vy ⓘ	Watson Machine Learning	Lite	Dallas	Not associated	Default

Cancel **Associate**



h. Go to **Home** using drop-down option in the left panel.

Please be aware that due to specific geographic and business constraints, some may not be eligible to register. If you encounter this limitation, you can still explore the vast landscape of Large Language Models through alternative services such as ChatGPT.

10. To experiment with a language model, click the first card **with Prompt Lab**, as shown in the screenshot below.

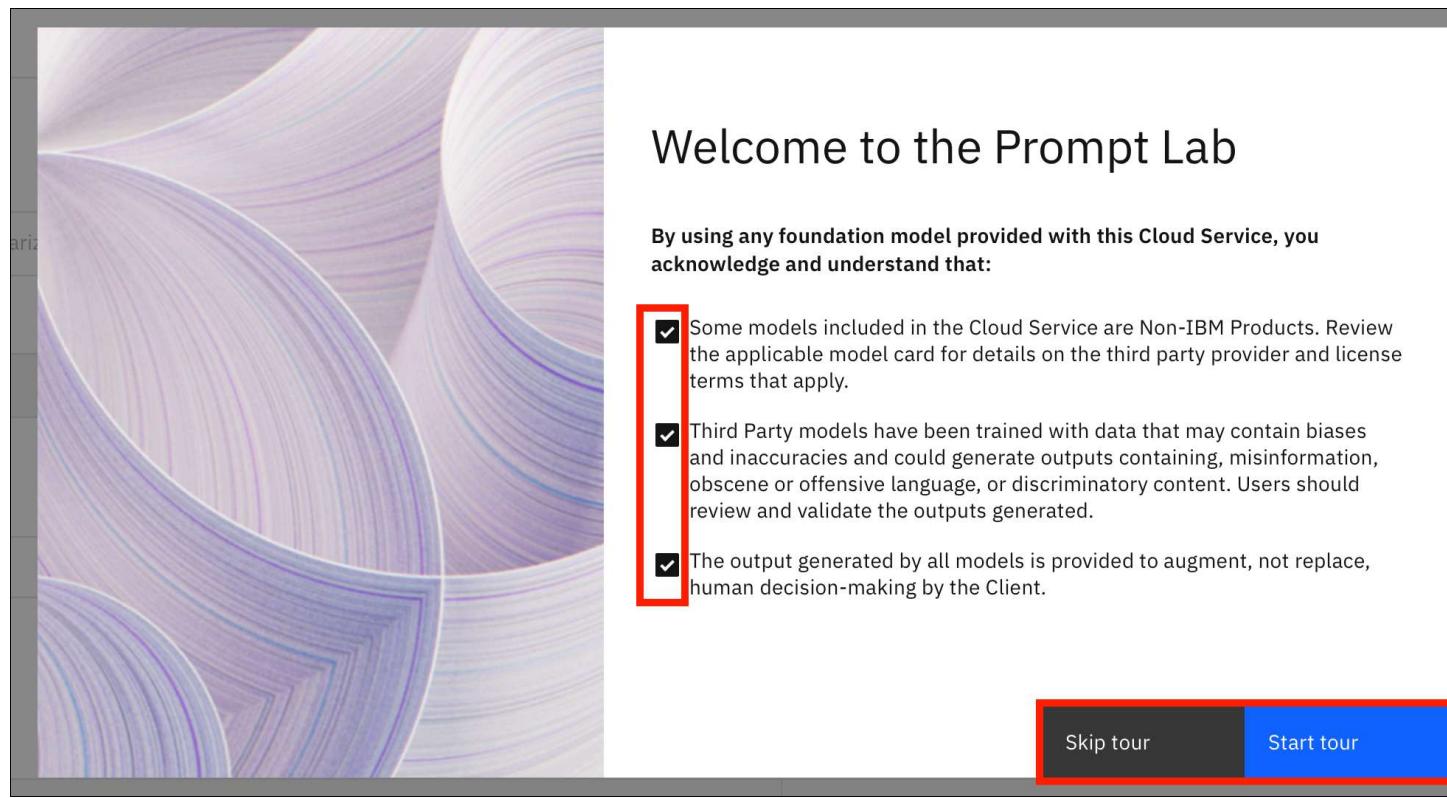


Note: If you have created a new project, click **Open Prompt Lab** in the Start working section

11. You will view the **Welcome to the Prompt Lab** tour page, as shown below.

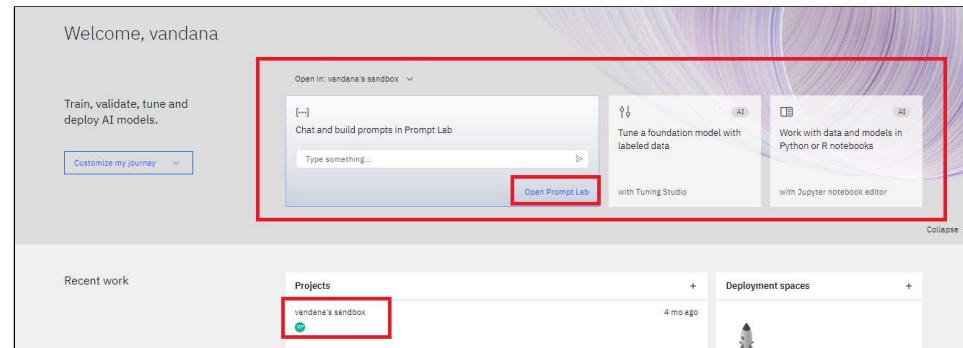
A screenshot of the 'Welcome to the Prompt Lab' tour page. On the left is a large, abstract, circular graphic composed of many thin, curved lines in shades of purple and blue. To the right of this graphic, the title 'Welcome to the Prompt Lab' is displayed in a large, bold, black font. Below the title, a block of text reads: 'By using any foundation model provided with this Cloud Service, you acknowledge and understand that:' followed by three bullet points. Each bullet point is preceded by an empty checkbox. At the bottom right of the page are two buttons: 'Skip tour' and 'Start tour'.

12. You must check all three acknowledgment boxes to start or skip the tour, as shown below.



Getting Started with watsonx

The Home page should look similar to this screen.



If you want to work with the Prompt Lab, you will need to work in the context of a project. If you look carefully, you will notice a highlighted square at the bottom of this screen. That's your personal sandbox. In this instance, it's labeled "Vandana's sandbox", as Vandana has signed in to her watsonx account.

Within each project, a suite of tools awaits, enabling you to interact with the AI, construct machine learning models, utilize Jupyter notebooks, and more. Of course, you're welcome to create more projects to organize and structure your work.

If you wish to do that, choose a project from the watsonx.ai home page, and click **Open Prompt Lab**. This will open the Prompt Lab.

Part A - Exploring the Prompt Lab

Upon entering the Prompt Lab within your project, you'll be greeted by this interface. Select **Structured** tab.

Familiarize yourself with the tool's features. On the left, you'll find icons representing Sample Prompts, Saved Prompts, and History. By default, these are collapsed.

Additionally, observe the options to create a New Prompt, save your current work, and turn on or off the AI Guardrails. These guardrails are essential safety measures designed to prevent the AI from producing harmful or offensive content. Feel free to explore these features.

Once you've had a good look around, let's delve deeper into the tool's primary components.

At the top, you'll spot three tabs: **Chat**, **Structured** and **Freeform**. As the name suggests, the Structured tab offers a guided approach to input, reminiscent of a more advanced version of the AI Classroom you've previously used in this course. On the other hand, the Freeform tab presents a simple textbox, allowing you to directly query the AI.

Take time to familiarize yourself with the Structured mode, as it offers a more guided experience.

There are three sections for you in here:

1. **Instruction Section:** This is akin to our earlier "Prompt Instructions."
2. **Example Section:** Here, you can provide specific input-output pairs to guide the AI, similar to our exercises in the Chain-of-Thought lab.
3. **Try Section:** This is where you'll input your queries. The AI's responses will be displayed in the Output field within this section.

Finally, note the **Generate** button located at the bottom right. Initially, it appears grayed out, indicating that no input has been provided. Once you enter your query, the button will become active, allowing you to submit your input to the AI for processing.

Exercise 1: Getting Started with Prompt Lab

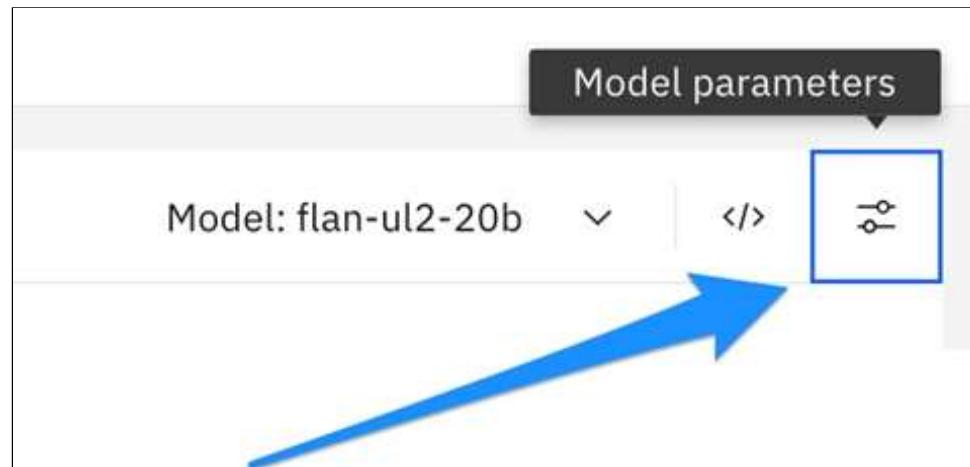
The following exercises should familiarize you with the Prompt Lab interface. For future convenience, consider saving prompts which you find particularly useful. You can use them as templates, simply tweaking the input field to suit your current query. This method is both efficient and time-saving. Select model **flan-ul2-20b**

1. **Basic Query:** Without any prompt instructions or examples, ask the AI:

Give me a list of common fish species in North America.

You might find the output lacking. For instance, I received `List of fish species in North America`. Where's the list? This truncation occurs in part because the default chat setup limits the response to a set word count. Let's rectify this.

2. **Adjusting Model Parameters:** Click the Model parameters settings icon in the top right, as depicted below.



Adjust the Max token parameter within a range of 1 to 1024. For this exercise, set it to 200, capping our output at 200 words. You can also set a minimum word count and tweak [more advanced parameters](#).

Remember, the [cost of using foundation models in IBM WatsonX](#) is determined by usage, which relates to the number of tokens generated. Therefore, setting a conservative Max tokens value is cost-effective.

Model parameters

Decoding

Greedy Sampling ⓘ

Repetition penalty

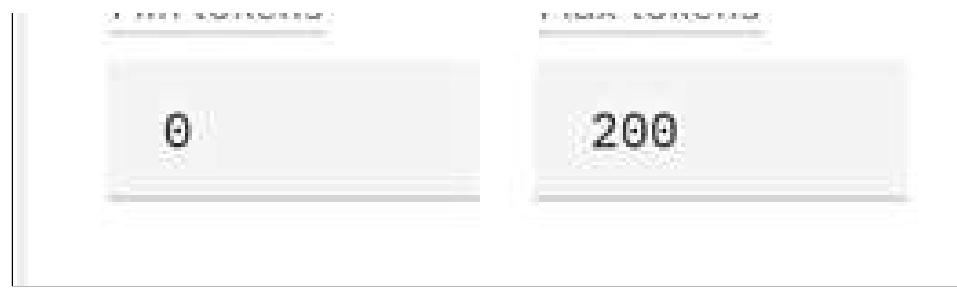
1 2 1

Stopping criteria ⓘ

Stop sequences

Min tokens

Max tokens



3. Switching Models: From the dropdown adjacent to Model parameters, switch from the `flan-ul2` model to IBM's `mixtral-8x7b-instruct-v01-q` or `granite-13b-chat-v2` model. Click Generate. You should now receive a comprehensive list of common fish species in North America. It doesn't apply here, but it's worth noting that if you're pleased with the output, you can also use the top right corner of the Output field to transfer the output to the Example section. This ensures future responses are modeled after this successful example.

4. Guardrail Safeguards: You might observe that one entry was omitted due to Guardrail safeguards. Temporarily disable the guardrail at the page's top and generate a new output. Can you identify which panfish species triggered the censor?

5. Improving English Clarity: Modify the instructions to state:

Acting as an English professor, you'll take my input, enhance its clarity, and make it sound more native.

Then, input the following fractured English sentence:

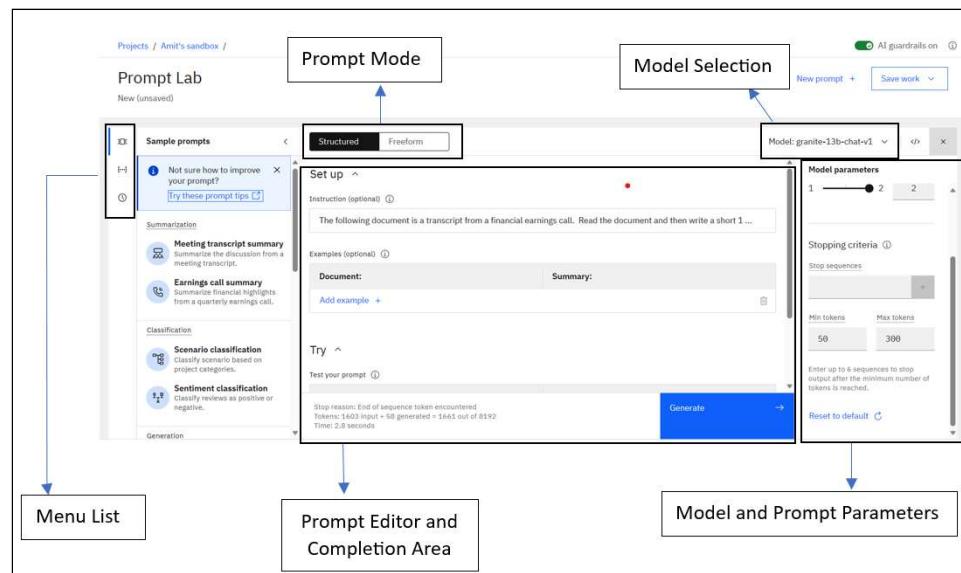
Yesterday, I go to big market for buy many fruits, but they no have the apple I like, so me feel sad and come back home without buy anything.

Experiment with various models. The results should be impressive, making this tool invaluable for non-native English speakers.

6. Translation: Use the prompt instructions to guide the model to translate your input into Italian or a second language you might speak. Then, input an English sentence. If you do speak a second language, try it with a few different models to see which one performs best. In my case, LLama by Meta did a particularly impressive job with Italian. watsonx is also smart enough to recommend when a specific prompt can benefit from an example, guiding you toward better results.

Part B - Prompting a foundation model

We can now explore in further detail, a few of the functions available on the watsonx Prompt Lab page with a focus on querying foundation models.



Menu List

The menu lists Sample prompts, which you can use. All you need to do is select the task-specific scenario, and the corresponding prompt will appear in Prompt Editor. This is a good way to learn how to prompt a foundation model.

Instructions to use:



Click on sample prompt icon.

Sample prompt may be expanded on the left side of Prompt Editor.

The screenshot shows the Microsoft AI Prompt Lab interface. On the left, there's a sidebar titled "Sample prompts" with several categories: "Meeting transcript summary", "Earnings call summary", "Classification", and "Sentiment classification". The "Meeting transcript summary" category is currently selected. In the main workspace, there's a "Set up" section with an "Interaction (optional)" input field containing the placeholder "Write a short summary for the meeting transcripts." Below it, there's a "Transcript" section with a list of messages from a meeting, and a "Summary" section where the AI has generated a summary based on the transcript. To the right, there are "Model parameters" for decoding (Greedy Sampling) and repetition penalty, and "Stopping criteria" for stop sequences. At the bottom, there's a "Generate" button and some token count information.

If you click on **Meeting transcript summary**, you will see the sample prompts listed for summarizing data from meeting transcripts.

The screenshot shows a modal window titled "Summarization". Inside, there's a sub-section titled "Meeting transcript summary" with the sub-instruction: "Summarize the discussion from a meeting transcript." There's also a small icon of three people at a table.

Under **Classification**, you will find sample prompts to help you sort data based on your preferences.

Model Selection

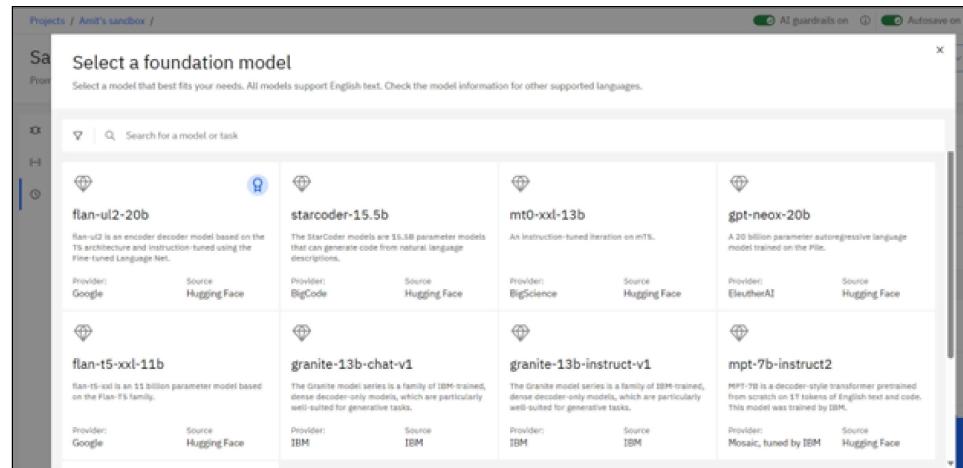
Another interesting feature is the Model Selection button.
Click on the down arrow to open menu.



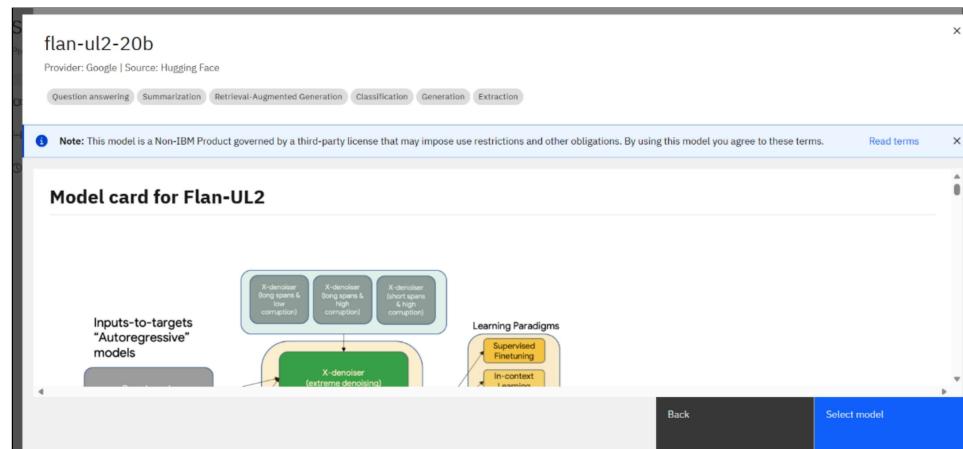
It will list some foundation models and allow you to explore more.

A screenshot of a sidebar menu titled "Recents". The sidebar has a white background with a thin gray border. It lists three items: "flan-ul2-20b", "granite-13b-chat-v1", and "llama-2-70b-chat". Below these items is a button labeled "View all foundation models" with a right-pointing arrow. A blue rectangular box highlights this button. On the right side of the sidebar, there is a vertical scroll bar with a dark gray track and a white slider. At the top of the sidebar, there is a header with the text "Model: flan-ul2-20b" and a small upward-pointing triangle icon. To the right of the header are two small icons: a double-angle bracket and a circular arrow.

Click on **View all foundation models**. This will open the **Select foundation model** dialog box.



Click on any model. This will open a model card or **Model Datasheet** that contains detailed information about the model.



To select a model, click on the **Select Model** button at the bottom right corner of the dialog box.

The model that you have selected will appear in the **Model Selection** box.



Model and Prompt Parameters

The Model and Prompt Parameters function allows you to specify parameters to control how the model generates output in response to your prompt. Here, you will find a list of parameters you can control in the Prompt Lab.

Decoding

Decoding is the process a model uses to choose the tokens in the generated output. Greedy decoding selects the token with the highest probability at each step of the decoding process. Sampling decoding is more variable and random, which is desirable in creative use cases. However, with greater variability comes the risk of nonsensical output.

Repetition Penalty

If you notice the result generated for your chosen prompt, model, and parameters consistently contains repetitive text, you can try adding a repetition penalty. The higher the penalty, the less likely the result will include repeated text.

Stopping Criteria

You can affect the length of the output generated by the model in two ways: Specifying stop sequences and setting Min tokens and Max tokens. Min tokens must be less than or equal to Max tokens.

Exercise 2: How to prompt a foundation model

Scenario

You work with a cloud software company, and a customer-support chatbot has received a message from a customer. You are tasked with classifying the customer's message as a question or a problem description so the chat can be routed to the correct support team. Let's see how you can adjust the model parameters to achieve this task.

Step 1

Click on the **New Prompt** button

Step 2

Select a model: either mt0-xxl-13b, flan-t5-xxl-11b, or flan-ul2-20b

The model must only return one of the class names, not be creative and make up new classes.

Step 3

Select the decoding as Greedy.

Step 4

The Repetition penalty is set to 1 by default.

Step 5

Select the Stopping criteria. Specify two stop sequences: "Question" and "Problem". After the model generates either of those words, it should stop.

Step 6

Set the Max tokens parameter to 5.

Model: flan-ul2-20b ▾ | </> X

Model parameters

Greedy Sampling ⓘ

Repetition penalty

1 2

Stopping criteria ⓘ

Stop sequences

+

Question X Problem X

Min tokens

Max tokens

→

Step 7

Craft the Prompt text: paste this zero-shot prompt text into the freeform prompt editor in Prompt Lab.

```
*“Classify this customer message into one of two classes: Question, Problem.
Class name: Question*
*Description: The customer is asking a technical question or a how-to question about our products or services.*
*Class name: Problem*
*Description: The customer is describing a problem they are having. They might say they are trying something, but it's not working. They might say they are getting an error or unexpected results.*
*Message: I'm having trouble registering for a new account.*
*Class name: “*
```

The screenshot shows the 'Prompt Lab' application window. At the top, there are buttons for 'New prompt' and 'Save work'. Below that, a dropdown menu shows 'Model: flan-ul2-20b'. The main area is titled 'I-I' and contains a 'Freeform' tab selected. Inside, there is a text input field with the following content:

```
Classify this customer message into one of two classes: Question, Problem.

Class name: Question
Description: The customer is asking a technical question or a how-to question about our products or services.

Class name: Problem
Description: The customer is describing a problem they are having. They might say they are trying something, but it's not working. They might say they are getting an error or unexpected results.

Message: I'm having trouble registering for a new account.
Class name:
```

At the bottom right of the input field is a blue 'Generate' button with a right-pointing arrow.

Step 8

Click **Generate** to see the results.

It will be able to classify the message to corresponding Class either Question or Problem.

This screenshot shows the same 'Prompt Lab' interface after generating the response. The 'Generate' button has been clicked, and the output is visible at the bottom of the input field:

```
Stop reason: Stop sequence encountered
Tokens: 107 Input + 1 generated = 108 out of 4096
Time: 2.3 seconds
```

Next to the input field is a blue 'Generate' button with a right-pointing arrow. To the right of the input field is a sidebar titled 'Model parameters' which includes 'Greedy' (radio button selected), 'Sampling' (radio button), 'Repetition penalty' (set to 1), and 'Stopping criteria' (with 'Question' and 'Problem' checked). At the bottom of the sidebar are 'Min tokens' (0) and 'Max tokens' (5).

Exercise 3: Generate a numbered list**Scenario**

Say, that you want to generate a numbered list on a given theme using a few-shot prompt.

Step 1

Click on the **New Prompt** Button.

Step 2

Chose a model: flan-ul2-20b is a good choice.

Step 3

Set the decoding to **Sampling** and set the following parameters:

1. Temperature: 0.7
2. Top P: 1
3. Top K: 50

Step 4

Set the **Random seed** to 9045 to get a different output each time you click Generate. Specify a different value for the Random seed parameter or clear the parameter.

Model: flan-ul2-20b ▾ </> X

Model parameters

Decoding

Greedy Sampling ⓘ

Temperature

0 — 2 0.7

Top P (nucleus sampling)

0 — 1 1

Top K

1 — 100 50

Random seed

→ 9045



Step 5

Identify the **Stopping criteria**:

Model: flan-ul2-20b ▾ | </> X

Model parameters

Repetition penalty

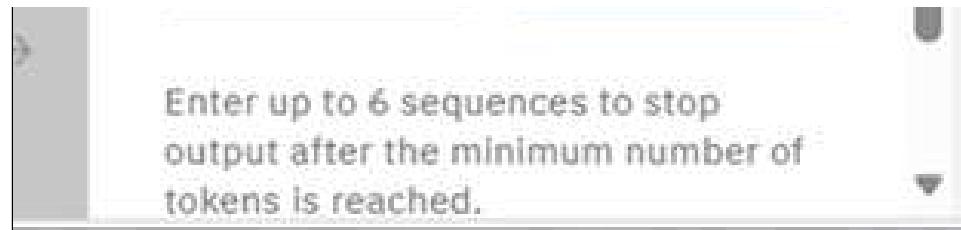
1 ● ————— 2

Stopping criteria ⓘ

Stop sequences

+
 ++ X

Min tokens Max tokens



To make sure the model stops generating text after one list, specify a stop sequence of two newline characters. To do that, click the **Stop sequence** text box, press the Enter key twice, then click Add sequence. Max tokens parameter to 50.

Step 6

Craft the prompt text: paste this few-shot prompt text into the freeform prompt editor in Prompt Lab.

```
*What are 4 types of dog breed?
1.Poodle
2.Dalmatian
3.Golden retriever
4.Bulldog*
*What are 3 ways to incorporate exercise into your day?
1.Go for a walk at lunch
2.Take the stairs instead of the elevator
3.Park farther away from your destination*
*What are 4 kinds of vegetable?
1.Spinach
2.Carrots
3.Broccoli
4.Cauliflower*
*What are the 3 primary colors?
1.Red
2.Green
3.Blue*
*What are 3 ingredients that are good on pizza?**
```

The screenshot shows the 'Prompt Lab' application interface. At the top, it says 'New (unsaved)' and has 'Save work' and 'X' buttons. Below that, the 'Model' dropdown is set to 'flan-uf2-20b'. On the left, there's a sidebar with 'Structured' and 'Freeform' tabs, and a list of recent files. The main area is titled 'Prompt Lab' and contains the following text:

```
I-I
2. Poodle
2. Dalmatian
3. Golden retriever
4. Bulldog*
What are 3 ways to incorporate exercise into your day?
1. Go for a walk at lunch
2. Take the stairs instead of the elevator
3. Park farther away from your destination*
What are 4 kinds of vegetable?
1. Spinach
2. Carrots
3. Broccoli
4. Cauliflower*
What are the 3 primary colors?
1. Red
2. Green
3. Blue*
What are 3 ingredients that are good on pizza?**
```

On the right side, there are 'Model parameters' (repetition penalty slider at 1), 'Stopping criteria' (stop sequences input field with a plus sign), and 'Min tokens' (0) and 'Max tokens' (50) input fields. A note at the bottom says: "Enter up to 6 sequences to stop output before the minimum number of tokens is reached."

Step 7

Click **Generate** to see the result.

The screenshot shows the 'Prompt Lab' application window. At the top, there are buttons for 'New prompt' and 'Save work'. The main area displays three separate AI-generated responses to different prompts:

- Q: What are 4 breeds of dog?**
A: 1. Poodle
2. Dalmatian
3. Golden retriever
4. Bulldog
- Q: What are 3 ways to incorporate exercise into your day?**
A: 1. Go for a walk at lunch
2. Take the stairs instead of the elevator
3. Park farther away from your destination
- Q: What are 4 kinds of vegetable?**
A: 1. Spinach
2. Carrots
3. Broccoli
4. Cauliflower

Below these, another prompt is shown:

Q: What are 3 primary colors?
A: 1. Red
2. Green
3. Blue

At the bottom, a final prompt is shown:

Q: What are 3 ingredients that are good on pizza?
A: 1. sausage 2. pepperoni 3. bacon

On the right side of the interface, there are 'Model parameters' settings for 'Repetition penalty' (set to 1), 'Stopping criteria', and 'Min tokens' (set to 0) and 'Max tokens' (set to 50). A note at the bottom right says 'Enter up to 6 sequences to stop output after the minimum number of tokens is reached.'

Three ingredients are suggested.

Conclusion

The Prompt Lab is a versatile tool. Its AI models have helpful capabilities such as generating creative content, answering complex queries, assisting in language translation, summarizing content, extracting specific data from unstructured data, and even simulating specific roles like an English professor, as seen in this lab. It's a testament to the power and utility of modern AI.



Skills Network